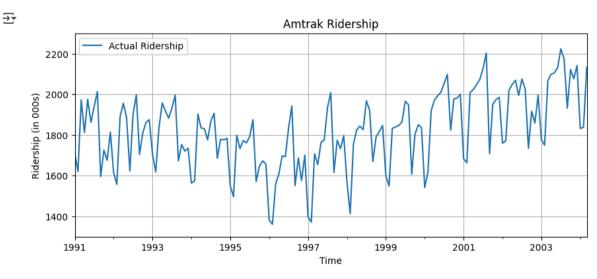
Amtrak Ridership Regression-Based Forecasting Model

Import Required Packages

```
# Importing libraries
import math
import numpy as np
import pandas as pd
import mathlotlib.pylab as plt
import statsmodels.formula.api as sm
import statsmodels.formula.api as smf
from statsmodels.tsa import tsatools, stattools
from statsmodels.tsa.arima_model import ARIMA
from statsmodels.graphics import tsaplots
```

Loading, Preprocessing and Splitting Data

```
# Load, convert Amtrak data for time series analysis
Amtrak_df = pd.read_csv('Amtrak.csv')
Amtrak_df.head(9)
print(Amtrak_df)
         Month
01/01/1991
                     Ridership
                      1708.917
         01/02/1991
                      1620.586
                      1972.715
         01/03/1991
         01/04/1991
                      1811.665
         01/05/1991
    154 01/11/2003
                      2076.054
    155
        01/12/2003
                      2140.677
         01/01/2004
                      1831.508
    156
                      1838.006
    157
         01/02/2004
    158 01/03/2004
                      2132.446
    [159 rows x 2 columns]
# Convert 'Month' to datetime and create time series
Amtrak_df['Date'] = pd.to_datetime(Amtrak_df['Month'], format='%d/%m/%Y')
ridership_ts = pd.Series(Amtrak_df['Ridership'].values, index=Amtrak_df['Date'])
# Add trend and constant to DataFrame
ridership_df = tsatools.add_trend(ridership_ts, trend='ct')
ridership_df['Ridership'] = ridership_ts.values # add the actual series for modeling
# Fit linear regression model: Ridership ~ trend + constant
ridership_lm = smf.ols(formula='Ridership ~ trend', data=ridership_df).fit()
# Plot the time series
ax = ridership_ts.plot(label='Actual Ridership', figsize=(10, 4))
ax.set_xlabel('Time')
ax.set_ylabel('Ridership (in 000s)')
ax.set_ylim(1300, 2300)
ax.set_title('Amtrak Ridership')
ax.legend()
plt.grid(True)
plt.show()
```



Data Splitting

```
# Split data: 120 training points, rest validation
train_df = ridership_df.iloc[:120]
valid_df = ridership_df.iloc[120:]
```

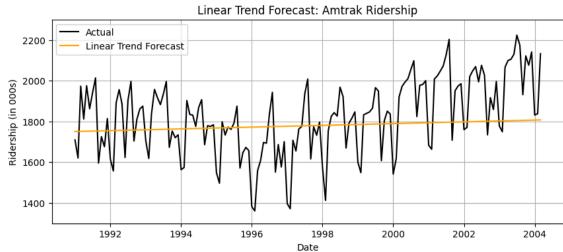
Linear Model

```
# Fit linear model on training data
ridership_lm = smf.ols(formula='Ridership ~ trend', data=train_df).fit()

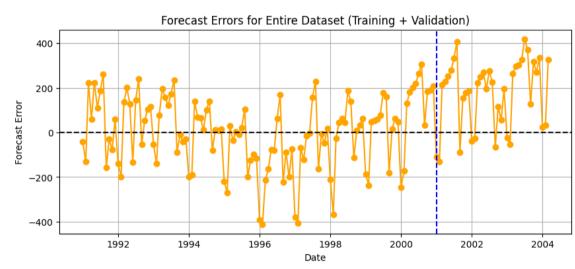
# Predict on full dataset
all_predictions = ridership_lm.predict(ridership_df)

# Plot actual vs linear trend
plt.figure(figsize=(10, 4))
plt.plot(ridership_df.index, ridership_df['Ridership'], label='Actual', color='black')
plt.plot(ridership_df.index, all_predictions, label='Linear Trend Forecast', color='orange')
plt.xlabel('Date')
plt.xlabel('Bidership (in 000s)')
plt.title('Linear Trend Forecast: Amtrak Ridership')
plt.ylim(1300, 2300)
plt.legend()
plt.grid(True)
plt.show()
```





```
# Compute forecast errors (actual - predicted)
all_errors = ridership_df['Ridership'] - all_predictions
# Plot forecast errors
plt.figure(figsize=(10, 4))
plt.plot(ridership_df.index, all_errors, color='orange', marker='o', linestyle='-')
plt.axhline(y=0, color='black', linestyle='--')
plt.xlabel('Date')
plt.ylabel('Forecast Error')
plt.title('Forecast Errors for Entire Dataset (Training + Validation)')
plt.grid(True)
# Mark the training/validation cutoff
cutoff_date = ridership_df.index[nTrain]
plt.axvline(x=cutoff_date, color='blue', linestyle='--', label='Training/Validation Split')
plt.show()
# View summary
print(ridership_lm.summary())
```



OLS Regression Results						
Dep. Variate Model: Method: Date: Time: No. Observator Model: Covariance	ations: Ls:	Least Squ Thu, 10 Jul 01:5	OLS Adj ares F-s 2025 Pro 9:38 Log 120 AIC 118 BIC	-	ic):	0.006 -0.002 0.7258 0.396 -778.50 1561.
	coe		t	P> t	[0.025	0.975]
Intercept trend	1750.054 0.359		59.426 0.852		1691.737 -0.477	1.196
Omnibus: Prob(Omnibus) Skew: Kurtosis:	ıs):	0 -0	.101 Jar .464 Pro	bin-Watson: que-Bera (JE b(JB):		1.065 4.645 0.0980 140.

Notes

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Interpretation

◆ Intercept (1750.05)

This represents the estimated baseline ridership (in thousands) when the trend value is zero — essentially, the ridership at the starting point of the dataset.

It is highly statistically significant (p < 0.001), meaning this estimate is reliable.

◆ Trend Coefficient (0.36)

This indicates that for each unit increase in the time trend (typically each month), ridership is expected to increase by 0.36 thousand passengers (i.e., 360 passengers).

 $However, the \ p-value \ is \ 0.396, which \ is \ not \ statistically \ significant \ (above \ the \ 0.05 \ threshold).$

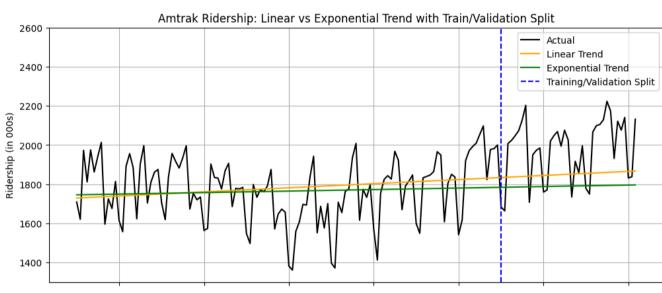
The 95% confidence interval [-0.477, 1.196] includes zero, further confirming that the trend effect is not significant.

✓ Conclusion:

While the baseline level of ridership is statistically strong, the trend over time is not significant, meaning there's no clear evidence of growth or decline in ridership based on this linear trend model.

Exponential Model

```
# Exponential model (log scale, then back-transform)
ridership_lm_expo = sm.ols(formula='np.log(Ridership) ~ trend', data=train_df).fit()
predict_log_expo_all = ridership_lm_expo.predict(ridership_df)
predict_expo_all = np.exp(predict_log_expo_all)
# Plot everything
plt.figure(figsize=(12, 5))
# Actual ridership
plt.plot(ridership_df.index, ridership_df['Ridership'], label='Actual', color='black')
# Linear trend forecast
plt.plot(ridership_df.index, predict_linear_all, label='Linear Trend', color='orange')
# Exponential trend forecast
plt.plot(ridership_df.index, predict_expo_all, label='Exponential Trend', color='green')
# Mark the training/validation cutoff
cutoff_date = ridership_df.index[nTrain]
plt.axvline(x=cutoff_date, color='blue', linestyle='--', label='Training/Validation Split')
# Final touches
plt.xlabel('Date')
plt.ylabel('Ridership (in 000s)')
plt.title('Amtrak Ridership: Linear vs Exponential Trend with Train/Validation Split')
plt.ylim(1300, 2600)
plt.legend()
plt.grid(True)
plt.show()
→
```



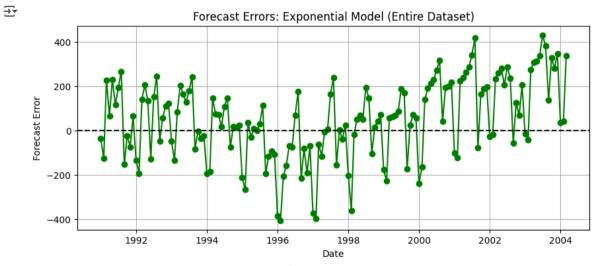
Calculate forecast errors for exponential model (actual - predicted)

forecast_errors_expo = ridership_df['Ridership'] - predict_expo_all

```
# Plot forecast errors
plt.figure(figsize=(10, 4))
plt.plot(ridership_df.index, forecast_errors_expo, color='green', marker='o', linestyle='-')
plt.axhline(y=0, color='black', linestyle='--')
plt.xlabel('Date')
plt.ylabel('Forecast Error')
plt.title('Forecast Errors: Exponential Model (Entire Dataset)')
plt.grid(True)
plt.show()
```

View summary

print(ridership_lm_expo.summary())



OLS Regression Results						
Dep. Variable:	np.	======================================	====== R–squar	ed:		0.005
Model:		0LS		-squared:		-0.004
Method:		Least Squares	F–stati			0.5448
Date:	Th	ı, 10 Jul 2025		-statistic):	0.462
Time:		01:59:58	Log-Lik	elihood:		115.58
No. Observations:		120	AIC:			-227.2
Df Residuals:		118	BIC:			-221.6
Df Model:		1				
Covariance Type:		nonrobust				
=======================================						========
(coef	std err	t	P>ltl	[0.025	0.9751

	coef	std err	t	P> t	[0.025	0.975]
Intercept trend	7.4646 0.0002	0.017 0.000	436.237 0.738	0.000 0.462	7.431 -0.000	7.498 0.001
Omnibus: Prob(Omnibus) Skew: Kurtosis:):	0. -0.		, -	:	1.066 9.314 0.00949 140.

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Interpretation

Intercept (7.4646)

This is the estimated log of the baseline ridership when the trend is zero.

Exponentiating the intercept gives the baseline ridership in the original scale:

[e^{7.4646} \approx 1746.8 \text{ thousand passengers}]

This estimate is highly statistically significant (p < 0.001), indicating a reliable baseline ridership level.

◆ Trend Coefficient (0.0002)

This coefficient represents the expected change in the log of ridership for each unit increase in the time trend (e.g., each month).

• Because the model is log-linear, the coefficient can be interpreted approximately as a **percentage change** in ridership per unit increase in trend:

[\text{Percentage change} \approx 0.0002 \times 100 = 0.02%]

- However, the p-value is 0.462, which is not statistically significant (greater than 0.05).
- The 95% confidence interval includes zero, indicating the trend effect is uncertain and not statistically meaningful.

Conclusion:

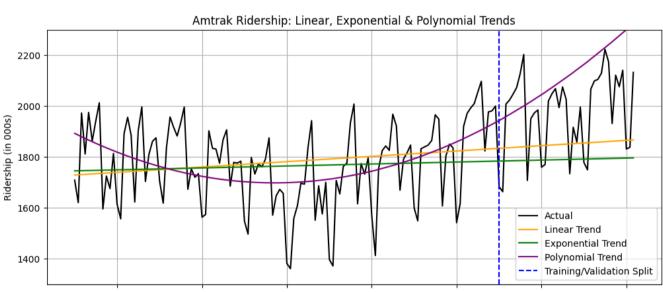
 $\overline{\mathbf{T}}$

The baseline ridership (on the original scale) is estimated reliably at about 1.75 million passengers.

However, there is no significant evidence of growth or decline in ridership over time based on the exponential (log-linear) model.

→ Polynomial Model

```
# Fit the polynomial (quadratic) trend model on training data
\label{eq:ridership_lm_poly} ridership\_lm\_poly = sm.ols(formula='Ridership \sim trend + np.square(trend)', data=train\_df).fit()
# Predict on full dataset using the polynomial model
# Add squared trend term to full data
ridership_df['trend_squared'] = np.square(ridership_df['trend'])
# Predict using polynomial model
predict_poly_all = ridership_lm_poly.predict(ridership_df)
# Plot all models including polynomial
plt.figure(figsize=(12, 5))
# Actual ridership
plt.plot(ridership_df.index, ridership_df['Ridership'], label='Actual', color='black')
plt.plot(ridership_df.index, predict_linear_all, label='Linear Trend', color='orange')
# Exponential trend
plt.plot(ridership_df.index, predict_expo_all, label='Exponential Trend', color='green')
plt.plot(ridership_df.index, predict_poly_all, label='Polynomial Trend', color='purple')
# Train-validation cutoff
cutoff_date = ridership_df.index[nTrain]
plt.axvline(x=cutoff_date, color='blue', linestyle='--', label='Training/Validation Split')
# Final touches
plt.xlabel('Date')
plt.ylabel('Ridership (in 000s)')
plt.title('Amtrak Ridership: Linear, Exponential & Polynomial Trends')
plt.ylim(1300, 2300)
plt.legend()
plt.grid(True)
plt.show()
```

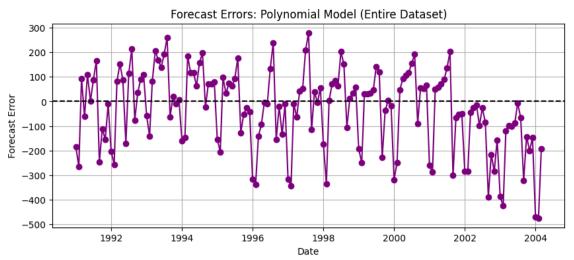


```
# Calculate forecast errors for polynomial model (actual - predicted)
forecast_errors_poly = ridership_df['Ridership'] - predict_poly_all

# Plot forecast errors
plt.figure(figsize=(10, 4))
plt.plot(ridership_df.index, forecast_errors_poly, color='purple', marker='o', linestyle='-')
plt.axhline(y=0, color='black', linestyle='--')
plt.xlabel('Date')
plt.ylabel('Forecast Error')
plt.title('Forecast Errors: Polynomial Model (Entire Dataset)')
plt.grid(True)
plt.show()

# View summary
print(ridership_lm_poly.summary())
```





OLS Regression Results

			=======================================
Dep. Variable:	Ridership	R-squared:	0.174
Model:	0LS	Adj. R-squared:	0.160
Method:	Least Squares	F-statistic:	12.31
Date:	Thu, 10 Jul 2025	<pre>Prob (F-statistic):</pre>	1.40e-05
Time:	02:00:39	Log-Likelihood:	-767.41
No. Observations:	120	AIC:	1541.
Df Residuals:	117	BIC:	1549.
Df Model:	2		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept trend np.square(trend)	1899.7695 -7.0031 0.0609	40.872 1.559 0.012	46.481 -4.491 4.874	0.000 0.000 0.000	1818.824 -10.091 0.036	1980.715 -3.915 0.086
Omnibus: Prob(Omnibus): Skew: Kurtosis:		6.745 0.034 -0.564 2.654	Durbin-Watso Jarque-Bera Prob(JB): Cond. No.		6 0.	281 6.951 0310 0e+04

Notes

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 1.99e+04. This might indicate that there are

strong multicollinearity or other numerical problems.

Interpretation

◆ Intercept (1899.77)

This is the estimated ridership (in thousands) at the starting point (when trend = 0). It reflects the baseline level and is **highly statistically significant** (p < 0.001).

Trend Coefficient (-7.00)

This negative coefficient suggests that initially, ridership decreases with time.

The significant p-value (p < 0.001) confirms that this declining effect is statistically meaningful.

Quadratic Term (trend² = 0.0609)

The positive and statistically significant (p < 0.001) quadratic term suggests that the rate of decline slows down and eventually **reverses into growth**.

This curvature indicates a **U-shaped trend**, where ridership first drops, then rises over time.

Model Interpretation Summary:

- The **polynomial model fits better** than the linear or exponential models, with **R**² = **0.174**, indicating that about 17.4% of the variation in ridership is explained by the model.
- The trend over time is **nonlinear**: there's an initial drop in ridership followed by a recovery.

Seasonal Model

```
# Add constant trend (just intercept) — no time trend this time
ridership_df = tsatools.add_trend(ridership_ts, trend='c')
# Add 'Month' column for seasonality (categorical)
ridership_df['Month'] = ridership_df.index.month
# Add 'Ridership' values again to use in regression
ridership_df['Ridership'] = ridership_ts.values
# Split into train and validation sets
nTrain = 120
train_df = ridership_df.iloc[:nTrain]
valid_df = ridership_df.iloc[nTrain:]
# Fit seasonal model using dummy variables for Month
ridership_lm_season = sm.ols(formula='Ridership ~ C(Month)', data=train_df).fit()
# Predict on full dataset
predict_season_all = ridership_lm_season.predict(ridership_df)
# Plot all models including seasonal
plt.figure(figsize=(12, 5))
# Actual ridership
plt.plot(ridership_df.index, ridership_df['Ridership'], label='Actual', color='black')
plt.plot(ridership_df.index, predict_season_all, label='Seasonal Model', color='brown')
# Train-validation split line
plt.axvline(x=ridership_df.index[nTrain], color='blue', linestyle='--', label='Training/Validation Split')
# Labels and legend
plt.xlabel('Date')
plt.ylabel('Ridership (in 000s)')
plt.title('Amtrak Ridership: Seasonal Model')
plt.ylim(1300, 2300)
```

→*

```
Amtrak Ridership: Seasonal Model

2200

Actual Seasonal Model
--- Training/Validation Split

1800

1400
```

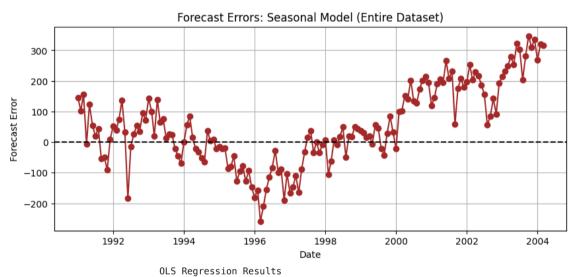
Compute forecast errors for all data
all_errors_season = ridership_df['Ridership'] - predict_season_all

Plot forecast errors
plt.figure(figsize=(10, 4))
plt.plot(ridership_df.index, all_errors_season, color='brown', marker='o', linestyle='-')
plt.axhline(y=0, color='black', linestyle='--')
plt.xlabel('Date')
plt.ylabel('Forecast Error')
plt.title('Forecast Errors: Seasonal Model (Entire Dataset)')
plt.grid(True)
plt.show()

View summary

print(ridership_lm_season.summary())

→



OLS Regression Results						
Dep. Variable:	Ridership	R-squared:	0.647			
Model:	0LS	Adj. R-squared:	0.611			
Method:	Least Squares	F-statistic:	18.01			
Date:	Thu, 10 Jul 2025	<pre>Prob (F-statistic):</pre>	8.13e-20			
Time:	02:01:25	Log-Likelihood:	-716.36			
No. Observations:	120	AIC:	1457.			
Df Residuals:	108	BIC:	1490.			
Df Model:	11					
Covariance Type:	nonrohust					

,,,,,,,,,,,,,,,,,,,,,,,,,,,,,,,,,,,,,,,						
	coef	std err	t	P> t	[0.025	0.975]
Intercept	1563.0546	31.568	49.514	0.000	1500.482	1625.627
C(Month)[T.2]	-45.2488	44.644	-1.014	0.313	-133.740	43.242
C(Month)[T.3]	254.3665	44.644	5.698	0.000	165.875	342.858
C(Month)[T.4]	256.0095	44.644	5.735	0.000	167.518	344.501
C(Month)[T.5]	289.1398	44.644	6.477	0.000	200.649	377.631
C(Month)[T.6]	244.3774	44.644	5.474	0.000	155.886	332.869
C(Month)[T.7]	356.2441	44.644	7.980	0.000	267.753	444.735
C(Month)[T.8]	407.5771	44.644	9.130	0.000	319.086	496.068
C(Month)[T.9]	86.6791	44.644	1.942	0.055	-1.812	175.170
C(Month)[T.10]	211.5252	44.644	4.738	0.000	123.034	300.016
C(Month)[T.11]	203.2728	44.644	4.553	0.000	114.782	291.764
C(Month)[T.12]	241.3327	44.644	5.406	0.000	152.841	329.824
Omnibus:		0.245	 Durbin-Wat	son:		0.398
Prob(Omnibus):		0 885	larque_Rer	a (IR).		0 267

Omnibus:	0.245	Durbin-Watson:	0.398
Prob(Omnibus):	0.885	Jarque-Bera (JB):	0.267
Skew:	-0.105	Prob(JB):	0.875
Kurtosis:	2.902	Cond. No.	12.9

Notes

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Interpretation

◆ Intercept (1563.05)

- $\bullet\,$ Baseline ridership (in thousands) for ${\bf January},$ when all dummy variables are zero.
- Highly significant (p < 0.001).

Monthly Effects (relative to January)

- March-August show large, significant increases (peaking in August at +407.58 k).
- **February** shows a non-significant decrease (-45.25 k, p = 0.313).
- **September** is marginal (p = 0.055) and may not differ from January.
- $\bullet \ \ \textbf{October-December} \ \text{all show significant positive differences}.$

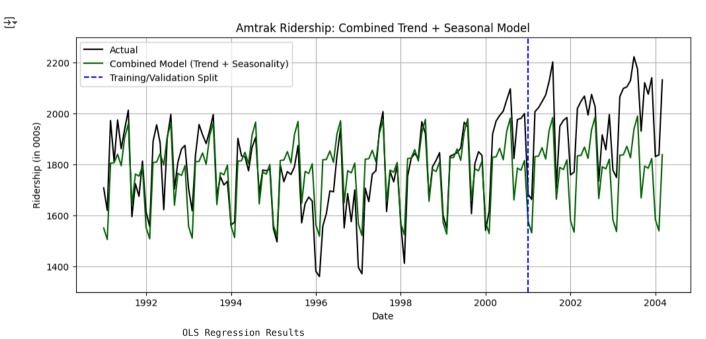
✓ Model Fit

• R² = 0.647, Adj. R² = 0.611: 64.7% of training variation explained.

- F-statistic = 18.01, p < 0.0001: Model is highly significant overall.
- Seasonal dummies capture strong month-to-month variation; adding a trend term (if desired) could further refine long-term direction.

Combined Trend and Seasonal Model

```
# Add linear trend column to ridership_df (if not already present)
# We'll create a numeric time index as the trend variable
ridership_df['trend'] = np.arange(len(ridership_df))
# Split train and validation sets again (to keep consistent)
train_df = ridership_df.iloc[:nTrain]
valid_df = ridership_df.iloc[nTrain:]
# Fit combined model: Ridership ~ trend + seasonality (Month)
\label{eq:ridership_lm_combined} \verb| sm.ols(formula='Ridership \sim trend + C(Month)', data=train_df).fit() \\
# Predict on the full dataset
predict_combined_all = ridership_lm_combined.predict(ridership_df)
# Plot actual ridership and combined model predictions
plt.figure(figsize=(12, 5))
plt.plot(ridership_df.index, ridership_df['Ridership'], label='Actual', color='black')
plt.plot(ridership_df.index, predict_combined_all, label='Combined Model (Trend + Seasonality)', color='darkgreen')
plt.xlabel('Date')
plt.ylabel('Ridership (in 000s)')
plt.title('Amtrak Ridership: Combined Trend + Seasonal Model')
plt.ylim(1300, 2300)
plt.legend()
plt.grid(True)
plt.show()
# View summary of combined model
print(ridership_lm_combined.summary())
```



===========			
Dep. Variable:	Ridership	R-squared:	0.649
Model:	0LS	Adj. R-squared:	0.610
Method:	Least Squares	F-statistic:	16.51
Date:	Thu, 10 Jul 2025	<pre>Prob (F-statistic):</pre>	2.60e-19
Time:	02:01:59	Log-Likelihood:	-715.99
No. Observations:	120	AIC:	1458.
Df Residuals:	107	BIC:	1494.
Df Model:	12		
Covariance Type:	nonrobust		

=======================================				=======		
	coef	std err	t	P> t	[0.025	0.975]
Intercept C(Month) [T.2] C(Month) [T.3] C(Month) [T.4] C(Month) [T.5] C(Month) [T.6] C(Month) [T.7]	1551.4031 -45.4646 253.9350 255.3622 288.2767 243.2986 354.9495	34.700 44.714 44.716 44.720 44.726 44.733 44.741	44.709 -1.017 5.679 5.710 6.445 5.439 7.933	0.000 0.312 0.000 0.000 0.000 0.000	1482.614 -134.105 165.290 166.710 199.613 154.621 266.255	1620.192 43.176 342.580 344.015 376.940 331.976 443.644
C(Month) [T.8] C(Month) [T.9] C(Month) [T.10] C(Month) [T.11] C(Month) [T.12] trend	406.0667 84.9530 209.5833 201.1151 238.9593 0.2158	44.752 44.763 44.777 44.791 44.808 0.265	9.074 1.898 4.681 4.490 5.333 0.815	0.000 0.060 0.000 0.000 0.000 0.417	317.352 -3.785 120.819 112.321 150.133 -0.309	494.781 173.691 298.348 289.909 327.786 0.741

Omnibus:	0.564	Durbin-Watson:	0.400			
Prob(Omnibus):	0.754	Jarque-Bera (JB):	0.642			
Skew:	-0.157	Prob(JB):	0.726			
Kurtosis:	2.827	Cond. No.	857.			

Interpretation

- Intercept (1551.40)
- Estimated ridership (in thousands) for **January** at trend = 0 (start of dataset).
- Highly statistically significant (p < 0.001).
- Monthly Effects (relative to January)
 - Months March to August show large, significant positive ridership increases (e.g., August +406k passengers).
 - February shows a non-significant decrease (\sim -45k, p=0.312).
 - September shows a positive but marginally non-significant increase (p=0.06).
 - October to December also show significant increases.
- Trend Coefficient (0.22)
 - Indicates an estimated increase of ~220 passengers per unit increase in time (likely monthly).

- However, the trend is **not statistically significant** (p = 0.417).
- The confidence interval [-0.31, 0.74] includes zero, so the model provides **no clear evidence of a linear upward or downward trend** beyond seasonality.

✓ Model Fit and Diagnostics

- R-squared = 0.649 and Adjusted R-squared = 0.610: About 61–65% of variation in ridership is explained by this model.
- **F-statistic = 16.51**, p-value << 0.001: Model is statistically significant overall.
- Durbin-Watson = 0.40: Indicates potential positive autocorrelation of residuals (values close to 2 are ideal).
- Omnibus and Jarque-Bera tests suggest residuals are approximately normally distributed.

⊀ Summary:

- The strong seasonal pattern dominates the variation in ridership.
- There is no significant linear trend detected after accounting for seasonality.
- Model effectively captures monthly ridership differences but may need refinement to address residual autocorrelation or explore non-linear trends.